databricks

GUIDE

The Agentic CDP RFP Guide

+ TEMPLATE

Is your CDP ready for the age of agentic buyers?



Contents

Introduction 3
Why the CDP Decision Has Changed 5
Four CDP Buying Scenarios 7
Defining an Agentic CDP 9
The Agentic CDP Litmus Test 11
Evaluation Framework on a Page 12
Requirement 1: Trusted Context and Identity 13
Requirement 2: Continuous 1:1 Learning 14
Requirement 3: Trust, Governance and Human Control 15
Requirement 4: Open Architecture and Interoperability 16
Requirement 5: Operating Model, Model Choice and Economics 17
Organizational Readiness 18
Architecture and Commercial Decision 20
Weighted Scorecard 22
Decision Paths 24
The Databricks Perspective 26

Introduction

Buying has entered the agentic era. Has your CDP?

Your customers now have AI agents working for them. Agents can search, compare and buy across more channels in seconds than any human could navigate in hours. Already, **74% of consumers have used AI tools to shop**, fundamentally changing the speed and complexity of the buying journey.

Brands are racing to respond. Gartner predicts that **60% of brands will use agentic AI to streamline 1:1 interactions by 2028**. Yet most CDP strategies were built for a traditional, campaign-based model focused on moving audiences through preplanned journeys.

That operating model cannot keep pace with signals and decisions moving at machine speed. Static segments and waterfall campaigns miss intent, produce generic experiences and leave agents without the customer, business and decision context required to act responsibly and in real time.

The underlying architecture is shifting, too. By 2030, Gartner predicts that **80% of net-new enterprise CDP deployments will be embedded in or composable** with data platforms rather than stand-alone. Leaders evaluating a CDP today are making a foundational decision about where customer context lives, how agents access it and how every decision remains governed and accountable.

74%

Of consumers have used AI tools to discover, compare or buy products

Source: CI&T

60%

Of brands will use agentic AI to streamline 1:1 interactions by 2028

Source: Gartner

80%

Of net-new enterprise CDP deployments are expected to be embedded in or composable with data platforms by 2030

Source: Gartner

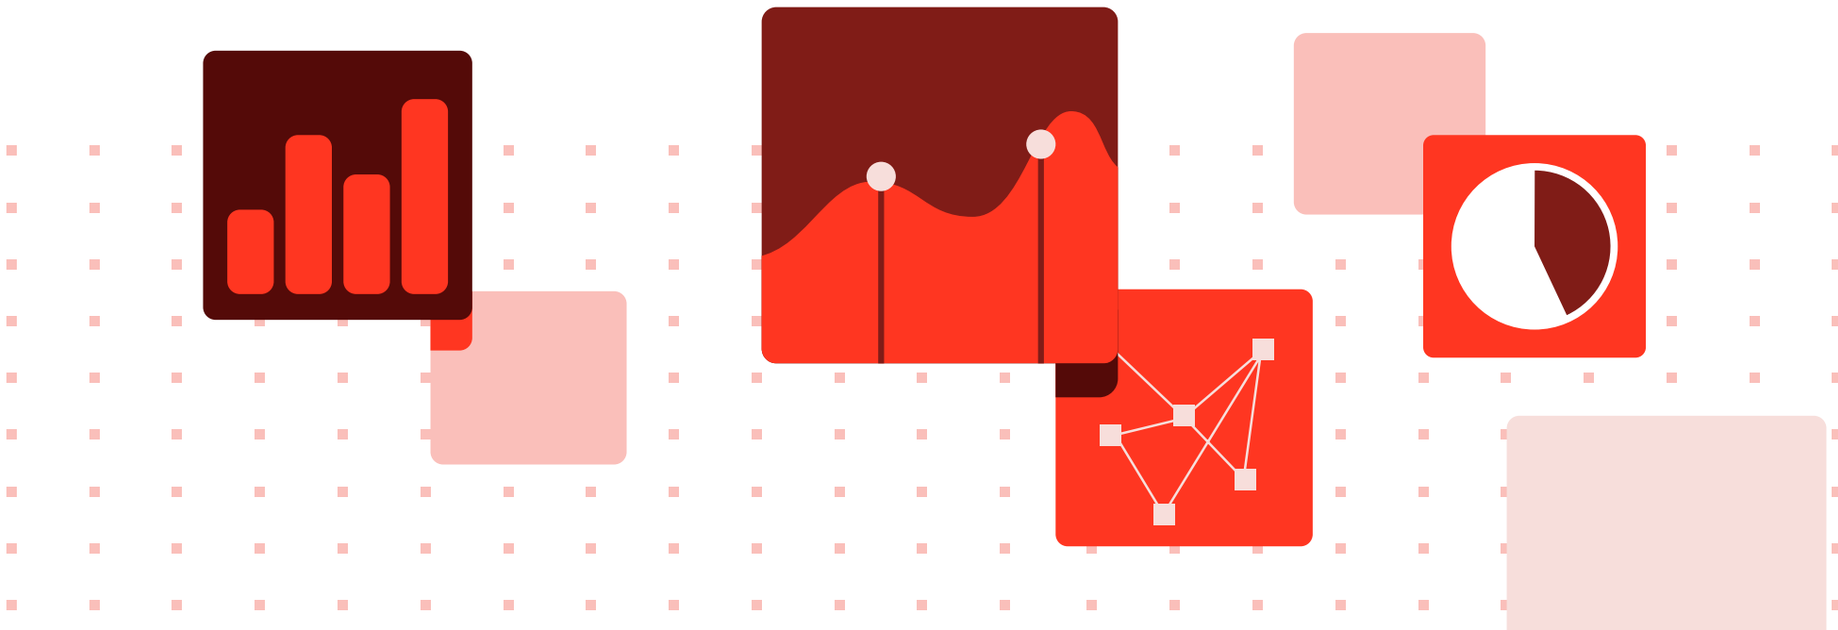
The decision

Choose a CDP that can turn customer data and context into continuous decisions and actions, without creating another data silo, locking the business into one ecosystem or weakening human accountability.

How to use this guide

This guide is a practical framework to help you align marketing, data and IT on CDP requirements, expose agent washing, test vendor claims and score platforms based on verified capabilities rather than feature lists or roadmap promises.

Use this guide to extend your existing CDP requirements with agentic evaluation criteria. Preserve your current requirements for data integration, identity resolution, segmentation, activation, measurement and governance, along with industry and business needs. Then apply the qualification questions, evaluation framework and scorecard to assess agentic capabilities. Use the results alongside your broader CDP evaluation when making a purchase decision.



Why the CDP Decision Has Changed

Buying moves at AI speed. Campaigns still move in waterfalls.

Buying is shifting from human-paced journeys to agent-mediated decisions. To keep up, marketing must move from static batch campaigns to continuous engagement loops that respond to real-time signals.

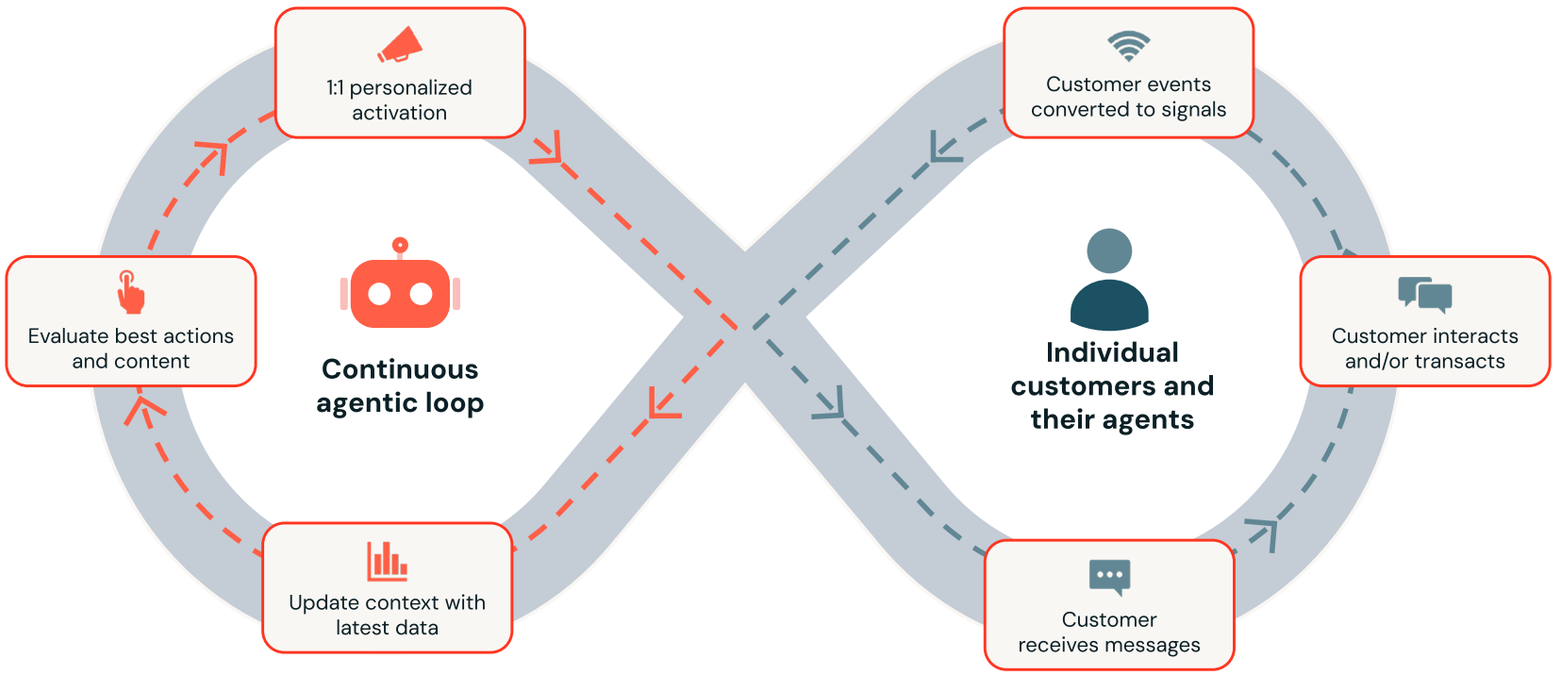
Yesterday's CDP was built for a slower buyer. Traditional CDPs were designed for sequential waterfall campaigns where teams create a plan, wait for data, build audiences and content, launch across channels and analyze performance after the fact. Every handoff adds delay. By the time the campaign reaches the customer, their intent, circumstances or needs may have already changed.

Waterfall Campaign Model



Agentic buying raises the stakes. AI compresses the cycle of discovery, comparison and action, meaning that brands need the ability to sense, decide and respond while the moment still matters.

To keep up in the agentic age, a CDP must be able to power continuous engagement loops that personalize every interaction, for every customer, in real time. Agents help marketers respond in the moment and learn from outcomes. Marketers bring the creative taste, brand judgment and understanding of their customers that define what a great experience should feel like. They set goals, shape experiences and determine the guardrails and approvals that guide agent action.



Continuous engagement loops that personalize every interaction

Four CDP Buying Scenarios

Four starting points. One higher standard.

Organizations arrive at an Agentic CDP evaluation with different assets, constraints and migration risks. Your starting point should determine what you preserve, replace and prove first. It should not change the standard for real-time context, continuous decisioning, governance and interoperability.

Starting point	What triggers the evaluation	What matters most	Common trap
Replace a packaged CDP or marketing suite	Renewal pressure, high TCO, slow audience delivery, duplicated data, suite lock-in or weak agent readiness	Migration continuity, identity quality, decommissioning cost, channel coverage, data portability and faster time to value	Running a feature-for-feature replacement that recreates the same silo on a newer interface
Evolve a Composable CDP	Customer data already lives in the data platform, while teams seek to simplify vendor complexity, align governance more closely, manage licensing costs or govern agentic capabilities	Which capabilities to retain, extend or consolidate, how the approach supports shared governance, the total cost across platforms and operations, and how new agents work with the existing activation layer	Assuming warehouse-native architecture automatically delivers unified governance, lower TCO or agentic readiness
Adopt a first CDP	Marketing needs self-service activation and agentic capabilities, while data and IT want to avoid creating another disconnected customer data estate	A measurable first use case, clear ownership, governed marketer access, time to value and an architecture that can grow	Buying platform breadth before agreeing on the first business outcome
Scale a DIY CDP	Internal teams have built profiles, models and activation pipelines, but maintenance demands are increasing or marketing still depends heavily on technical teams	Which capabilities and IP to retain, marketer experience, build-versus-buy economics, support requirements and portability	Ignoring internal labor and maintenance costs or buying software that duplicates mature internal capabilities

Start with one decision loop

Before issuing an RFP, define the first loop the platform must improve: the customer signal that starts it, the context required, the decision to make, the action permitted, the outcome to observe and the business metric that proves value.

The principle: Use your starting point to determine the migration path, then evaluate every option against the same Agentic CDP standard.



Defining an Agentic CDP

An Agentic CDP turns governed context into continuous action

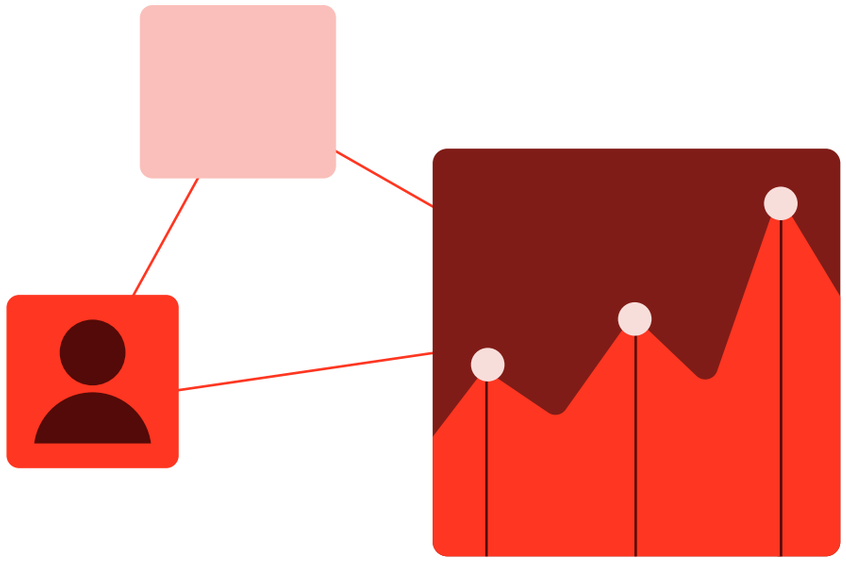
The category combines a complete CDP foundation with a new operating model in which agents analyze, decide, act and learn alongside humans.

An Agentic CDP is a customer data platform purpose-built for AI agents to operate alongside marketers. It gives agents governed customer, business and decision context so they can pursue goals, make bounded decisions, take action across channels, learn from outcomes and improve continuously. Humans define the strategy, permissions, guardrails and approval thresholds.

The CDP foundation

Use the agentic readiness criteria in this guide alongside your foundational CDP requirements. The following capabilities remain essential:

- Connect to customer data across relevant sources
- Resolve identities and maintain accurate profiles
- Build segments and audiences
- Activate through existing engagement channels
- Measure outcomes and return signals
- Govern access, consent, quality and lineage



The agentic operating model

An Agentic CDP must be:

- **Context aware:** Understands customer signals, business goals, current constraints, prior decisions and previous campaign performance
- **Goal directed:** Pursues a defined outcome within human-set objectives and boundaries
- **Autonomous:** Plans and completes multistep work across systems, including making decisions and executing engagement actions
- **Continuous:** Incorporates new signals and outcomes into each customer's context so learning informs subsequent decisions across campaigns
- **Accountable:** Makes every decision and action observable, explainable, auditable and subject to human approval or override

AI-enabled features don't make a CDP agentic

Adding a chatbot, natural-language audience builder or AI recommendations to a traditional CDP leaves the same human-operated, campaign-based workflow in place. These technologies can make a CDP easier to use, but they don't make a CDP truly agentic.

An Agentic CDP changes the operating model. A traditional CDP crosses the threshold to being an agentic CDP when agents can use timely, governed context to decide, act and learn continuously, without a marketer manually restarting each stage.

The next step is determining whether a platform can actually enable an agentic operating model.

The Agentic CDP Litmus Test

Five questions that expose agent washing

Use this test before advancing a vendor into detailed evaluation. Ask vendors to demonstrate all five capabilities using a relevant customer use case. Each answer should be supported by working functionality available today — roadmap commitments don't count as a pass.

#	Test	Pass/fail question
1	**Context aware**	Can agents use governed customer, business and decision context, including constraints and prior outcomes?
2	**Goal directed**	Can agents pursue a defined business goal and choose among possible next actions as conditions change, rather than following only a fixed workflow?
3	**Autonomous**	Can agents complete multistep engagement workflows, from segmentation and personalization to triggering sends, within defined permissions and approval rules?
4	**Continuous**	Can agents retain and use context and engagement outcomes for each customer to improve subsequent decisions without a marketer rebuilding or relaunching the campaign?
5	**Accountable**	Can teams set permissions, constraints, approval thresholds and stop conditions while tracing, evaluating, pausing and overriding every decision and action?

Evaluation Framework on a Page

Five requirements for an Agentic CDP

The Agentic CDP litmus test establishes whether a platform qualifies as agentic. The following weighted framework evaluates how well it can deliver in your environment.

When using this framework, score demonstrated capabilities using your data, use cases, constraints and existing ecosystem. All five areas are required. The weighting reflects their relative importance, not whether a requirement can be ignored.

Requirement area	Suggested weighting	What it establishes
Trusted context and identity	30%	Agents operate on accurate, current and governed customer, business and decision context grounded in resolved identities and reliable data.
Continuous 1:1 learning	25%	Agents use business goals and customer context to execute engagement actions, learn from outcomes and carry that learning into subsequent interactions and campaigns.
Trust, governance and human control	20%	Every access, decision and action is secure, explainable, observable, auditable and subject to human approval or override.
Open architecture and interoperability	15%	The platform works across the existing data, AI and martech ecosystem without unnecessary data copies or forced dependence on one vendor.
Operating model, model choice and economics	10%	Marketing gains autonomy while data and IT retain control, with flexible model options, transparent costs and contractual flexibility.
Total	100%	This structure provides a balanced framework for comparing platforms that have passed the Agentic CDP litmus test.

REQUIREMENT 1

Trusted Context and Identity

Build context agents can trust

A unified customer profile is the starting point. Agents also need current business conditions, consent and identity confidence, along with the history of decisions and outcomes. That context must remain governed as it changes.

REQUIRED CAPABILITIES

- Connect or federate transactional, behavioral, product, service, consent and operational data through batch and streaming methods
- Resolve identities across people, accounts, households, devices and interactions using deterministic, probabilistic and supervised methods
- Expose source lineage, freshness, consent status and data-quality signals
- Incorporate business context such as eligibility, inventory, margin, risk, policies and channel constraints
- Retain customer interaction history, prior decisions and outcomes as governed context that authorized agents can use across campaigns
- Correct, merge or split identities and propagate changes, restrictions and deletion requests downstream

QUESTIONS TO ASK

- Which data remains in our existing environment, which is federated, and which is copied into the platform?
- How quickly do new signals, consent changes and business conditions become available to an agent?
- How does the platform resolve conflicting identifiers and represent uncertain matches?
- Can a marketer or data steward inspect why two records were linked?
- How are consent and purpose-of-use policies carried from profile creation through decisioning and activation?
- Can we retain our existing identity rules, data models, identity partners and matching logic?

REQUIREMENT 2

Continuous 1:1 Learning

Turn every signal into a better next decision

The CDP must connect each customer's signals, decisions, actions and outcomes so that learning informs future engagement. It should respond in real time where the use case requires it, while updating decisions as reliable outcome evidence becomes available. Response speed and learning cadence should be evaluated separately.

REQUIRED CAPABILITIES

- Define an objective, success measure, eligible actions, constraints and time horizon for each engagement loop
- Use shared customer context and applicable goals, consent, eligibility and contact limits to inform engagement decisions
- Select an action, offer, message, channel, timing or suppression decision at the individual level
- Execute actions through existing engagement systems and confirm delivery or completion
- Connect outcomes to prior decisions and actions, then use that evidence to improve subsequent engagement for each customer

QUESTIONS TO ASK

- Which parts of the loop can run continuously, and where is human intervention still required?
- How do agents access and use a customer's prior interactions and outcomes when making decisions in a new campaign?
- Can you demonstrate how an observed outcome changes a customer's next action, offer, channel, timing or suppression decision?
- Can the platform optimize toward a business outcome rather than a proxy such as clicks or opens?

REQUIREMENT 3

Trust, Governance and Human Control

Keep every agent decision inside a control boundary

Agents increase the speed and scale of both value and risk. Governance must apply while the agent is operating and across the data it can access, the decisions it can make, the tools it can use and the actions it can take.

REQUIRED CAPABILITIES

- Apply role-based and least-privilege access by user, agent, data type, action, channel, market and business unit
- Enforce consent, purpose, retention and customer-contact policies at decision time
- Record the context, objective, rationale, policy version, action and outcome for each decision
- Provide visibility into agent activity and a way for authorized users to stop or intervene when needed
- Pause, override, roll back or safely redirect an agent when conditions or risks change

QUESTIONS TO ASK

- How are conflicting policies resolved, and which policy takes precedence?
- What records and explanations are available to investigate a decision or action?
- Which actions require approval, and who can approve them?
- How does the platform detect unauthorized data access, tool use, actions or unexpected agent behavior?
- How long are decision and action records retained, and can they be exported for audit or investigation?

REQUIREMENT 4

Open Architecture and Interoperability

Keep the customer stack open and adaptable

An Agentic CDP should work across the organization's existing data, model and engagement ecosystem. Openness requires production-ready, bidirectional interfaces and a practical path to change components without rebuilding the foundation.

REQUIRED CAPABILITIES

- Connect to existing engagement systems through documented interfaces that support activation and return outcome signals
- Use customer context from existing data platforms while minimizing unnecessary movement and duplication
- Support external models, identity providers, measurement tools, content systems and custom business logic
- Keep customer data, profiles, decision history and business logic portable, and support channel, model or other component changes without rebuilding the customer data foundation

QUESTIONS TO ASK

- Which integrations are production-ready and bidirectional, and what service levels apply?
- What data must leave our governed environment, where is it copied, how long is it retained and what does movement cost?
- Can we use our own models, identity logic, business rules and agent tools?
- What changes if we replace an activation destination, identity provider or model vendor?
- What data, configurations, policies and decision history can we export at contract end?

REQUIREMENT 5

Operating Model, Model Choice and Economics

Build an operating model that can scale

The platform must fit how marketing, data, IT, risk, procurement and finance work together. Marketers need governed autonomy, technical teams need model and architecture choices, and leaders need a clear view of cost and value.

REQUIRED CAPABILITIES

- Provide role-based workflows that let marketers operate within reusable data, policy and approval controls
- Separate strategy, policy definition, execution, approval and oversight responsibilities
- Support development, testing, simulation, versioning, approval and controlled production release
- Use commercial, open or custom models and route tasks based on quality, latency, cost and risk
- Attribute model, compute, data, activation and service costs to markets, teams and outcomes
- Support budgets, fallbacks, phased deployment and commercially practical exit terms

QUESTIONS TO ASK

- What can marketers change directly, and which changes require data, IT or risk approval?
- Who owns each stage of configuring, approving, deploying, monitoring and improving an engagement loop?
- Can we use our own models and evaluation sets, and can the model for a task be changed without rebuilding the workflow?
- How does the platform compare models across output quality, latency, risk and cost?
- How does pricing scale with users, profiles, data volume, decisions, model calls, connectors, egress and services?
- Can we reduce scope, change providers or exit while retaining our data, logic, models and decision history?

Organizational Readiness

Determine autonomy by use case

Platform capability and organizational readiness must advance together. Before an agent can act for the brand, teams must define the outcome, provide trusted context, set enforceable boundaries and assign accountability. Readiness determines where autonomy can begin and what evidence is required to expand it.

Five readiness gates

Outcome: A measurable business outcome, accountable owner, baseline and success threshold are defined

Ownership: Marketing, data, IT and risk responsibilities are clear, including who approves, monitors, intervenes and answers for results

Context: The required customer data, metrics and business priorities are accessible, current, governed and reliable enough to support the decision

Feedback: A meaningful outcome signal returns quickly and consistently enough to evaluate performance and improve the next decision

Control: Consent, brand, legal, risk and budget policies are explicit, enforceable and supported by escalation, pause and recovery procedures

A practical autonomy ladder

Level 0: Observe — The agent analyzes activity and explains findings. Humans decide and act

Level 1: Recommend — The agent proposes a decision or action. Humans review and execute

Level 2: Execute after approval — The agent prepares the work and acts only after an authorized person approves it

Level 3: Execute within limits — The agent takes permitted, reversible actions within defined policies and escalates exceptions

Level 4: Adapt within limits — The agent continuously adjusts decisions based on outcomes while humans monitor performance, investigate exceptions and retain override authority

Match human control to risk

As Scott Brinker writes in Not all humans-in-the-loop are equal: One useful way to decide how much responsibility to introduce — and where — is to consider uncertainty, consequence and reversibility.

- **Uncertainty:** How much judgment does the decision require?
- **Consequence:** What is the impact if the decision is wrong?
- **Reversibility:** How quickly and completely can the action be undone?

Greater uncertainty requires more review, greater consequence requires deeper expertise and greater reversibility allows humans to operate farther from each individual action.

Recommended starting point: Begin with a narrow, measurable and reversible use case at Level 1 or 2. Expand autonomy only after demonstrating reliable data, evaluation coverage, controlled economics and effective incident response.

The principle: Assess readiness for each use case. A team may automate refill reminders while maintaining direct approval for pricing, credit or regulated communications.

Architecture and Commercial Decision

Make control a contract requirement

Architecture determines where data, governance and switching costs accumulate. Commercial terms determine whether those costs remain visible and controllable as agent usage grows. Evaluate both against the future state you intend to operate.

Architecture requirements

- **Map data residency:** Document where profiles, context, decision history, evaluations and model outputs physically reside and who controls them
- **Minimize duplication:** Require a business and technical justification for every copy of customer data, including its retention, protection, synchronization and deletion
- **Extend governance:** Verify that permissions, lineage, consent and business semantics remain enforceable across models, agents and activation systems
- **Test resilience:** Define how the system responds when a model, data source, policy service or channel becomes unavailable
- **Prove replaceability:** Demonstrate that models, identity services, channels and other components can change without rebuilding the customer data foundation
- **Require portability:** Establish how data, logic, policies, evaluations and decision history can be exported and reused in accessible formats

Commercial requirements

- **Expose every cost driver:** Document how data, compute, storage, model usage, decisions, profiles, seats, connectors, egress, support and services affect price
- **Price realistic workloads:** Model pilot, expected and high-growth scenarios using documented assumptions for volume, latency, model mix and decision frequency
- **Account for existing investments:** Identify which data, governance, compute and AI capabilities the proposal reuses and which it duplicates
- **Make usage controllable:** Require cost attribution by use case, team and model, with budgets, alerts, rate limits and the ability to reduce or stop consumption
- **Create a practical ramp:** Align pilot terms, capacity commitments and pricing tiers with demonstrated adoption and value
- **Protect future leverage:** Define renewal limits, price-change protections, termination rights, export timelines, transition support and ownership of developed assets

Decision principles

Choose an architecture that keeps data, governance and intellectual property under your control, supported by commercial terms that make every cost visible, attributable and adjustable.

- **Reuse before replicating** — Favor architectures that build on existing governed data, AI and infrastructure investments
- **Price the future state** — Forecast continuous, individual-level decisions across use cases rather than relying on today's profile or audience volumes
- **Preserve choice** — Favor open formats, portable context, replaceable models, transparent consumption and commitments that expand with proven value

Weighted Scorecard

Turn evidence into a decision

Use one scorecard for each vendor and priority use case. Score each area from 0 to 5, then convert that score into weighted points.

Download the editable RFP worksheet here.

Scoring standard

0 — Absent, claim-only or roadmap-only

1 — Available in a narrow form with major capability gaps

2 — Partially meets the requirement or depends on manual workarounds

3 — Meets the requirement in a buyer-defined test

4 — Exceeds the requirement in a realistic, end-to-end test

5 — Differentiated and proven at production-scale with credible evidence



The following scorecard evaluates agentic readiness. Use it alongside your foundational CDP, industry and business requirements.

Requirement area	Weight	What a score of 3 must establish	Score	Weighted points
Trusted context and identity	**30%**	Agents can use accurate, current and governed customer and business context, including identity confidence, consent, constraints and prior outcomes.	**__/5**	**__/30**
Continuous 1:1 learning	**25%**	The platform demonstrates an individual engagement decision and permitted action, then shows how the outcome updates customer context and informs a subsequent decision.	**__/5**	**__/25**
Trust, governance and human control	**20%**	Access, decisions and actions are policy-bound, traceable, auditable, monitored and subject to approval, intervention or override.	**__/5**	**__/20**
Open architecture and interoperability	**15%**	The platform works bidirectionally across the existing ecosystem, minimizes unnecessary data copies and supports component replacement and export.	**__/5**	**__/15**
Operating model, model choice and economics	**10%**	Teams gain governed autonomy, can choose and route models, attribute and control costs, reuse existing investments and retain practical exit rights.	**__/5**	**__/10**
Total	**100%**	All five areas are required.		**__/100**

To calculate a vendor's total score

1. For each requirement area, use the following formula to determine weighted points:
   (score ÷ 5) × weight = weighted points
   For example, for a score of 4 in an area weighted at 30%:
   (4 ÷ 5) × 30 = 24 weighted points
2. Add the weighted points from all five areas to calculate the vendor's total score out of 100.

Decision Paths

Apply the gates before choosing a path

The weighted score is one input to the decision. An Agentic CDP vendor must qualify as genuinely agentic, prove the priority use case and fit your organization's readiness, architecture and commercial requirements.

Four decision gates

1. **Agentic qualification**
   Does the vendor pass all five questions in the Agentic CDP litmus test using capabilities available today?
2. **Required capabilities**
   Does the vendor score at least 3 in every requirement area based on sufficient evidence?
3. **Priority use case**
   Has the vendor demonstrated the complete loop using buyer-defined signals, context, policies, actions and success measures?
4. **Operational viability**
   Can the organization govern the proposed autonomy, and are the architecture, economics, service levels and contract terms acceptable?

Decision rule: Remove a vendor from the shortlist if it fails gate 1 or gate 2. A high total can't compensate for a required capability gap.

Choose the appropriate path

PROCEED	MODIFY
All four gates are satisfied, every area scores at least 3, the weighted total is 80 or higher and no critical risks remain unresolved. **NEXT STEP** Advance to commercial and implementation planning. Put use-case outcomes, cost limits and governance controls into the agreement.	The vendor clears the required capability gates, but the total is 60 to 79, or a manageable gap remains in scope, autonomy, architecture, economics or contract terms. **NEXT STEP** Narrow the initial use case, phase autonomy, define corrective milestones and retest before expanding scope or commitment.
PAUSE	**PRESERVE OPTIONALITY**
A required area scores below 3, critical proof remains only on the roadmap or the organization can't yet govern the proposed autonomy. **NEXT STEP** Address the internal prerequisite, maintain negotiating flexibility and resume evaluation against a defined readiness milestone.	No vendor clears the required standard, important capabilities remain roadmap-only or the current platform can support near-term needs while the category matures. **NEXT STEP** Seek a shorter renewal, secure export rights, avoid new proprietary dependencies and use a bounded pilot to gather evidence.

Decision principle: The highest-scoring vendor still has to clear every required gate. Don't reward a vendor for being the least incomplete option. If none qualify, preserving flexibility is a valid strategic outcome.

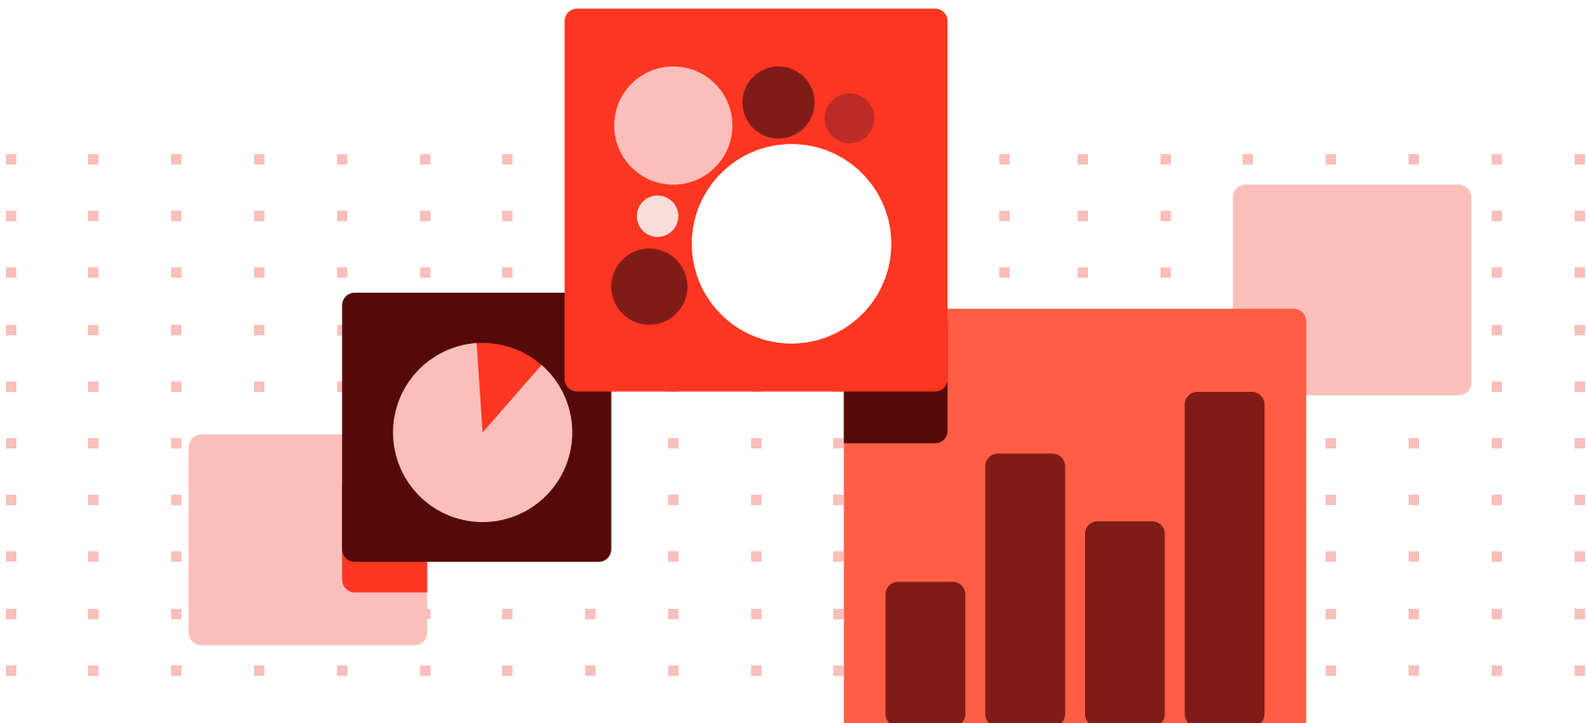
The Databricks Perspective

The CDP belongs in your data and AI foundation

Our point of view is simple: AI agents should operate where governed customer context, models and enterprise controls already live. Bringing CDP capabilities to that foundation reduces duplication, preserves governance and is the only way to enable agents to make relevant and accurate decisions at scale.

Three principles for the Agentic CDP

- **Embedded** — Build customer profiles, intelligence and activation on the governed data and AI foundation rather than creating another customer data store to secure and reconcile
- **Democratized** — Give marketers purpose-built, agent-first workflows while data and IT retain control of trusted data, models, policies and access
- **Autonomous** — Progress from one-off campaigns to infinity campaigns: continuous decision loops in which agents analyze, decide, act and learn within human-defined goals and guardrails



Introducing Databricks CustomerLake

CustomerLake is the Agentic CDP built natively in the lakehouse. It's designed to bring Customer 360, identity resolution, segmentation, activation and personalization into the same platform where enterprises already manage governed data, AI and business context.

- **Profile Agents:** Help data teams prepare customer data, identify quality issues, support third-party enrichment and unify fragmented records into business-ready Customer 360 profiles
- **Agentic Identity Resolution (AIR):** Combines deterministic, probabilistic and agentic workflows, with LLM and human review for edge cases and a feedback loop that improves matching logic over time
- **Campaign Agents:** Help marketers build audiences, recommend next-best actions, activate through existing channels and continuously improve engagement around business goals via Infinity Campaigns
- **An open ecosystem:** Using Reverse ETL, OpenSharing and Lakehouse Federation preserves existing martech and adtech investments while reducing unnecessary data movement, duplicated infrastructure and disconnected governance
- **A value-aligned consumption model:** Lets organizations reuse existing data, governance and AI investments while scaling spend with adoption and usage

Hold every vendor to the same standard

Verify current availability, test each capability against your own use case and award scorecard credit only for evidence demonstrated today. Use the product direction to inform your architecture and RFP requirements while preserving flexibility around timing and commitment.

Set your requirements before the market sets them for you

Your next CDP decision will determine where customer intelligence lives, how agents use it and how much control your organization retains. The requirements you set today will shape your marketing operating model for years to come. Use this guide to demand evidence, preserve flexibility and choose a foundation that lets marketing move at the speed of agents without compromising governance or control.

Set a higher standard for your next CDP

The requirements you set today will shape how your brand engages customers for years to come. Use this guide to strengthen your RFP criteria, test vendor claims and choose a foundation that connects customer context, agent action and continuous learning.

Explore the Databricks vision for CustomerLake and a new marketing operating model built on governed data and AI.

Learn more about CustomerLake



About Databricks

Databricks is the data and AI company. More than 20,000 organizations worldwide — including adidas, AT&T, Bayer, Block, Mastercard, Rivian, Unilever and 70% of the Fortune 500 — rely on Databricks to build and scale data and AI apps, analytics and agents. Headquartered in San Francisco, with 30+ offices around the globe, Databricks offers a unified platform that includes Lakebase, Genie, Agent Bricks, Lakeflow, Lakehouse and Unity Catalog. To learn more, follow Databricks on LinkedIn, X, YouTube and Instagram.

databricks

© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark and the Spark logo are trademarks of the Apache Software Foundation. Privacy Policy | Terms of Use